

Predicting the prices of Avocados

About the data-

The dataset represents weekly 2018 retail scan data for National retail volume (units) and price. Retail scan data comes directly from retailers' cash registers based on actual retail sales of Hass avocados. Starting in 2013, the table below reflects an expanded, multi-outlet retail data set. Multi-outlet reporting includes an aggregation of the following channels: grocery, mass, club, drug, dollar and military. The Average Price (of avocados) in the table reflects a per unit (per avocado) cost, even when multiple units (avocados) are sold in bags. The Product Lookup codes (PLU's) in the table are only for Hass avocados. Other varieties of avocados (e.g. greenskins) are not included in this table.

Some relevant columns in the dataset:

- Date - The date of the observation
- AveragePrice - the average price of a single avocado
- type - conventional or organic
- year - the year
- Region - the city or region of the observation
- Total Volume - Total number of avocados sold
- 4046 - Total number of avocados with PLU 4046 sold
- 4225 - Total number of avocados with PLU 4225 sold
- 4770 - Total number of avocados with PLU 4770 sold

```
In [1]: # display image using python
from PIL import Image
from IPython.display import display
img = Image.open(r"C:\Users\ankus\Desktop\NareshIT\2. Notes\11.Machine learning\
display(img)
```



Importing Libraries

```
In [2]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
import warnings  
warnings.filterwarnings('ignore')
```

```
In [3]: # import dataset  
data = pd.read_csv(r"C:\Users\ankus\Desktop\NareshIT\2. Notes\11.Machine learning  
data.head()
```

Out[3]:	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bags	Sr B
0	0	2015-12-27	1.33	64236.62	1036.74	54454.85	48.16	8696.87	8603
1	1	2015-12-20	1.35	54876.98	674.28	44638.81	58.33	9505.56	9408
2	2	2015-12-13	0.93	118220.22	794.70	109149.67	130.50	8145.35	8042
3	3	2015-12-06	1.08	78992.15	1132.00	71976.41	72.58	5811.16	5671
4	4	2015-11-29	1.28	51039.60	941.48	43838.39	75.78	6183.95	5986


```
In [4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18249 entries, 0 to 18248
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Unnamed: 0    18249 non-null   int64  
 1   Date         18249 non-null   object  
 2   AveragePrice 18249 non-null   float64 
 3   Total Volume 18249 non-null   float64 
 4   4046        18249 non-null   float64 
 5   4225        18249 non-null   float64 
 6   4770        18249 non-null   float64 
 7   Total Bags   18249 non-null   float64 
 8   Small Bags   18249 non-null   float64 
 9   Large Bags   18249 non-null   float64 
 10  XLarge Bags  18249 non-null   float64 
 11  type         18249 non-null   object  
 12  year         18249 non-null   int64  
 13  region       18249 non-null   object  
dtypes: float64(9), int64(2), object(3)
memory usage: 1.9+ MB
```

There are three categorical features and luckily no missing value. Let's explore the data further.


```
In [5]: sns.distplot(data['AveragePrice'])
```

```
Out[5]: <Axes: xlabel='AveragePrice', ylabel='Density'>
```



```
In [6]: sns.countplot(x='year', data=data, hue='type')
```

```
Out[6]: <Axes: xlabel='year', ylabel='count'>
```




There are almost equal numbers of conventional and organic avacados.Though,there is very less observation in the year 2018

```
In [7]: data.year.value_counts()
```

```
Out[7]: year
2017    5722
2016    5616
2015    5615
2018     1296
Name: count, dtype: int64
```

```
In [8]: sns.boxplot(y='type',x='AveragePrice',data=data,hue='type')
```


```
Out[8]: <Axes: xlabel='AveragePrice', ylabel='type'>
```



Organic avocados are more expensive. This is obvious, because their cultivation is more expensive and we all love natural products and are willing to pay a higher price for them.

```
In [9]: data.year = data.year.apply(str)
sns.boxenplot(x='year', y='AveragePrice', data=data, hue='year')
```

```
Out[9]: <Axes: xlabel='year', ylabel='AveragePrice'>
```



Avacados were slightly more expensive in the year 2017.(as there was shortage due to some reasons)

Dealing with categorical features

```
In [10]: data['type'] = data['type'].map({'conventional':0,'organic':1})  
  
# Extracting month from date column  
data.Date = data.Date.apply(pd.to_datetime)  
data['Month'] = data['Date'].apply(lambda x:x.month)  
data.drop('Date',axis=1,inplace=True)  
data.Month = data.Month.map({1:'JAN',2:'FEB',3:'MARCH',4:'APRIL',5:'MAY',6:'JUNE'})
```

```
In [20]: plt.figure(figsize=(9,5))  
sns.countplot(x='count',y='Month',data=data)  
plt.title('Monthwise Distribution of Sales',fontdict = {'fontsize':25})
```

```
-----  
TypeError Traceback (most recent call last)  
Cell In[20], line 2  
      1 plt.figure(figsize=(9,5))  
----> 2 sns.countplot(x='count',y='Month',data=data)  
      3 plt.title('Monthwise Distribution of Sales',fontdict = {'fontsize':25})  
  
File ~\anaconda3\Lib\site-packages\seaborn\categorical.py:2629, in countplot(data, x, y, hue, order, hue_order, orient, color, palette, saturation, fill, hue_norm, stat, width, dodge, gap, log_scale, native_scale, formatter, legend, ax, **kwargs)  
    2627     y = 1 if list(x) else None  
    2628 elif x is not None and y is not None:  
-> 2629     raise TypeError("Cannot pass values for both `x` and `y`.")  
    2631 p = _CategoricalAggPlotter(  
    2632     data=data,  
    2633     variables=dict(x=x, y=y, hue=hue),  
    2634     (...)  
    2635     legend=legend,  
    2636     )  
    2637     2638 )  
    2639 if ax is None:  
  
TypeError: Cannot pass values for both `x` and `y`.  
<Figure size 900x500 with 0 Axes>
```

It implies that sales of avacado see a rise in January,February and March

Preparing data for ML models

```
In [12]: # creating dummy variables  
dummies = pd.get_dummies(data[['year','region','Month']],drop_first=True)  
data_dummies = pd.concat([data[['Total Volume','4046','4225','4770','Total Bags'  
    'Large Bags','XLarge Bags','type']],dummies],axis=1)  
target = data['AveragePrice']
```

```
In [13]: # Splitting data  
from sklearn.model_selection import train_test_split  
X_train,X_test,y_train,y_test = train_test_split(data_dummies,target,test_size=0
```

```
#Standardizing the data
cols_to_std = ['Total Volume', '4046', '4225', '4770', 'Total Bags', 'Small Bags'
from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
scaler.fit(X_train[cols_to_std])
X_train[cols_to_std] = scaler.transform(X_train[cols_to_std])
X_test[cols_to_std] = scaler.transform(X_test[cols_to_std])
```

importing ML models from scikit-learn

In [16]:

```
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.svm import SVR
from sklearn.neighbors import KNeighborsRegressor
from xgboost import XGBRegressor
from sklearn.metrics import mean_absolute_error,mean_squared_error,r2_score
```

In [18]:

```
#to save time all models can be applied once using for loop
regressors = {
    'Linear Regression' : LinearRegression(),
    'Decision Tree' : DecisionTreeRegressor(),
    'Random Forest' : RandomForestRegressor(),
    'Support Vector Machines' : SVR(gamma=1),
    'K-nearest Neighbors' : KNeighborsRegressor(n_neighbors=1),
    'XGBoost' : XGBRegressor()
}

results = pd.DataFrame(columns=['MAE','MSE','R2-score'])
for method,func in regressors.items():
    model = func.fit(X_train,y_train)
    pred = model.predict(X_test)
    results.loc[method] = [np.round(mean_absolute_error(y_test,pred),3),
                          np.round(mean_squared_error(y_test,pred),3),
                          np.round(r2_score(y_test,pred),3)]
]
results
```

Out[18]:

	MAE	MSE	R2-score
Linear Regression	0.180	0.058	0.635
Decision Tree	0.133	0.043	0.732
Random Forest	0.096	0.019	0.878
Support Vector Machines	0.116	0.027	0.830
K-nearest Neighbors	0.100	0.023	0.854
XGBoost	0.094	0.017	0.895

Conclusion:

- Except linear regression model, all other models have mean absolute error less than 10% of mean of target variable.

- For this dataset, XGBoost and Random Forest algorithms have shown best results.

In []: