

Import Libraries

```
In [2]: import numpy as np
import pandas as pd
%matplotlib inline
import matplotlib.pyplot as plt
```

Creating DataFrames

```
In [3]: list_of_dicts=[{"name":"Ginger","breed":"Dachshund","height_cm":22,"weight_kg":10,"date_of_birth":"2019-03-14"}, {"name":"Scout","breed":"Dalmatian","height_cm":59,"weight_kg":25,"date_of_birth":"2019-05-09"}]
new_dogs = pd.DataFrame(list_of_dicts)
new_dogs
```

```
Out[3]:   name      breed  height_cm  weight_kg  date_of_birth
0   Ginger    Dachshund        22         10  2019-03-14
1     Scout    Dalmatian        59         25  2019-05-09
```

```
In [4]: dict_of_lists ={ "name": ["Ginger", "Scout"], "breed": ["Dachshund", "Dalmatian"], "height_cm": [22, 59], "weight_kg": [10, 25], "date_of_birth": ["2019-03-14", "2019-05-09"] }
new_dogs = pd.DataFrame(dict_of_lists)
new_dogs
```

```
Out[4]:   name      breed  height_cm  weight_kg  date_of_birth
0   Ginger    Dachshund        22         10  2019-03-14
1     Scout    Dalmatian        59         25  2019-05-09
```

Reading and Writing CSV files

```
In [5]: avocado = pd.read_csv(r"C:\Users\ankus\Desktop\NareshIT\2. Notes\11.Machine learning\avocado.csv")
avocado.head() # first 5 rows returns
```

Out[5]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bags	Sr B
0	0	2015-12-27	1.33	64236.62	1036.74	54454.85	48.16	8696.87	8603.62
1	1	2015-12-20	1.35	54876.98	674.28	44638.81	58.33	9505.56	9408.07
2	2	2015-12-13	0.93	118220.22	794.70	109149.67	130.50	8145.35	8042.21
3	3	2015-12-06	1.08	78992.15	1132.00	71976.41	72.58	5811.16	5677.40
4	4	2015-11-29	1.28	51039.60	941.48	43838.39	75.78	6183.95	5986.26

Read CSV and assign index

You can assign columns as index using "index_col" attribute.

Since I want to index Date there is another helpful function called "parse_date" which will parse the date in the rows such that we can perform more complex subsetting(eg monthly, weekly etc)

In [6]:

```
avocado = pd.read_csv(r"C:\Users\ankus\Desktop\NareshIT\2. Notes\11.Machine learning\avocado.csv")
avocado.head()
```

Out[6]:

	Unnamed: 0	AveragePrice	Total Volume	4046	4225	4770	Total Bags	Small Bags
	Date							
2015-12-27	0	1.33	64236.62	1036.74	54454.85	48.16	8696.87	8603.62
2015-12-20	1	1.35	54876.98	674.28	44638.81	58.33	9505.56	9408.07
2015-12-13	2	0.93	118220.22	794.70	109149.67	130.50	8145.35	8042.21
2015-12-06	3	1.08	78992.15	1132.00	71976.41	72.58	5811.16	5677.40
2015-11-29	4	1.28	51039.60	941.48	43838.39	75.78	6183.95	5986.26

Remove index from dataframe.reset_index(drop)

```
In [7]: avocado = avocado.reset_index(drop=True)
avocado.head()
```

Out[7]:

	Unnamed: 0	AveragePrice	Total Volume	4046	4225	4770	Total Bags	Small Bags	Large Bags
0	0	1.33	64236.62	1036.74	54454.85	48.16	8696.87	8603.62	9
1	1	1.35	54876.98	674.28	44638.81	58.33	9505.56	9408.07	9
2	2	0.93	118220.22	794.70	109149.67	130.50	8145.35	8042.21	10
3	3	1.08	78992.15	1132.00	71976.41	72.58	5811.16	5677.40	13
4	4	1.28	51039.60	941.48	43838.39	75.78	6183.95	5986.26	19

To write a CSV file function dataframe.to_csv(FILE_NAME)

```
In [8]: avocado.to_csv("test_write.csv")
```

Some useful pandas function

```
In [9]: avocado = pd.read_csv(r"C:\Users\ankus\Desktop\NareshIT\2. Notes\11.Machine learning\avocado.csv")
avocado.head()
```

Out[9]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bags	Small Bags	Large Bags
0	0	2015-12-27	1.33	64236.62	1036.74	54454.85	48.16	8696.87	8603.62	9
1	1	2015-12-20	1.35	54876.98	674.28	44638.81	58.33	9505.56	9408.07	9
2	2	2015-12-13	0.93	118220.22	794.70	109149.67	130.50	8145.35	8042.21	10
3	3	2015-12-06	1.08	78992.15	1132.00	71976.41	72.58	5811.16	5677.40	13
4	4	2015-11-29	1.28	51039.60	941.48	43838.39	75.78	6183.95	5986.26	19

```
In [10]: avocado.tail()
```

Out[10]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bags
18244	7	2018-02-04	1.63	17074.83	2046.96	1529.20	0.00	13498.67
18245	8	2018-01-28	1.71	13888.04	1191.70	3431.50	0.00	9264.84
18246	9	2018-01-21	1.87	13766.76	1191.92	2452.79	727.94	9394.11
18247	10	2018-01-14	1.93	16205.22	1527.63	2981.04	727.01	10969.54
18248	11	2018-01-07	1.62	17489.58	2894.77	2356.13	224.53	12014.15

In [11]: `avocado.info()` *#get concise summary of dataframe*

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18249 entries, 0 to 18248
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        18249 non-null   int64  
 1   Date              18249 non-null   object 
 2   AveragePrice      18249 non-null   float64
 3   Total Volume     18249 non-null   float64
 4   4046              18249 non-null   float64
 5   4225              18249 non-null   float64
 6   4770              18249 non-null   float64
 7   Total Bags        18249 non-null   float64
 8   Small Bags        18249 non-null   float64
 9   Large Bags        18249 non-null   float64
 10  XLarge Bags       18249 non-null   float64
 11  type              18249 non-null   object 
 12  year              18249 non-null   int64  
 13  region             18249 non-null   object 
dtypes: float64(9), int64(2), object(3)
memory usage: 1.9+ MB
```

In [12]: `avocado.shape`

Out[12]: (18249, 14)

In [13]: `avocado.describe()` *#statistical description*

Out[13]:

	Unnamed: 0	AveragePrice	Total Volume	4046	4225	
count	18249.000000	18249.000000	1.824900e+04	1.824900e+04	1.824900e+04	1.824900
mean	24.232232	1.405978	8.506440e+05	2.930084e+05	2.951546e+05	2.283974
std	15.481045	0.402677	3.453545e+06	1.264989e+06	1.204120e+06	1.074641
min	0.000000	0.440000	8.456000e+01	0.000000e+00	0.000000e+00	0.000000
25%	10.000000	1.100000	1.083858e+04	8.540700e+02	3.008780e+03	0.000000
50%	24.000000	1.370000	1.073768e+05	8.645300e+03	2.906102e+04	1.849900
75%	38.000000	1.660000	4.329623e+05	1.110202e+05	1.502069e+05	6.243420
max	52.000000	3.250000	6.250565e+07	2.274362e+07	2.047057e+07	2.546439

◀ ▶

In [14]: `avocado.values` *#numpy representation of given dataframe*

Out[14]: `array([[0, '2015-12-27', 1.33, ..., 'conventional', 2015, 'Albany'], [1, '2015-12-20', 1.35, ..., 'conventional', 2015, 'Albany'], [2, '2015-12-13', 0.93, ..., 'conventional', 2015, 'Albany'], ..., [9, '2018-01-21', 1.87, ..., 'organic', 2018, 'WestTexNewMexico'], [10, '2018-01-14', 1.93, ..., 'organic', 2018, 'WestTexNewMexico'], [11, '2018-01-07', 1.62, ..., 'organic', 2018, 'WestTexNewMexico']], dtype=object)`

In [15]: `avocado.columns`

Out[15]: `Index(['Unnamed: 0', 'Date', 'AveragePrice', 'Total Volume', '4046', '4225', '4770', 'Total Bags', 'Small Bags', 'Large Bags', 'XLarge Bags', 'type', 'year', 'region'], dtype='object')`

In [16]: `len(avocado.columns)`

Out[16]: 14

Sorting

In [17]: `# sort values based on "AveragePrice" (ascending) and "year" (descending)
avocado.sort_values(["AveragePrice", "year"], ascending = [True, False])`

Out[17]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	
15261	43	2017-03-05	0.44	64057.04	223.84	4748.88	0.00	!
7412	47	2017-02-05	0.46	2200550.27	1200632.86	531226.65	18324.93	45
15473	43	2017-03-05	0.48	50890.73	717.57	4138.84	0.00	4
15262	44	2017-02-26	0.49	44024.03	252.79	4472.68	0.00	3
1716	0	2015-12-27	0.49	1137707.43	738314.80	286858.37	11642.46	10
...
16720	18	2017-08-27	3.04	12656.32	419.06	4851.90	145.09	
16055	42	2017-03-12	3.05	2068.26	1043.83	77.36	0.00	
14124	7	2016-11-06	3.12	19043.80	5898.49	10039.34	0.00	
17428	37	2017-04-16	3.17	3018.56	1255.55	82.31	0.00	
14125	8	2016-10-30	3.25	16700.94	2325.93	11142.85	0.00	

18249 rows × 14 columns



Subsetting

Subsetting is used to get a slice of the original dataframe

In [20]:

```
# Subsetting columns
avocado["AveragePrice"]
```

Out[20]:

```
0      1.33
1      1.35
2      0.93
3      1.08
4      1.28
...
18244    1.63
18245    1.71
18246    1.87
18247    1.93
18248    1.62
Name: AveragePrice, Length: 18249, dtype: float64
```

Subsetting multiple columns

```
In [22]: avocado[["AveragePrice", "Date"]]
```

```
Out[22]:
```

	AveragePrice	Date
0	1.33	2015-12-27
1	1.35	2015-12-20
2	0.93	2015-12-13
3	1.08	2015-12-06
4	1.28	2015-11-29
...
18244	1.63	2018-02-04
18245	1.71	2018-01-28
18246	1.87	2018-01-21
18247	1.93	2018-01-14
18248	1.62	2018-01-07

18249 rows × 2 columns

Subsetting rows

```
In [29]: avocado["AveragePrice"] < 1
```

```
Out[29]:
```

0	False
1	False
2	True
3	False
4	False
...	
18244	False
18245	False
18246	False
18247	False
18248	False

Name: AveragePrice, Length: 18249, dtype: bool

```
In [24]: #This will print only the rows with price < 1  
avocado[avocado["AveragePrice"] < 1]
```

Out[24]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Buy
2	2	2015-12-13	0.93	118220.22	794.70	109149.67	130.50	8145.
6	6	2015-11-15	0.99	83453.76	1368.92	73672.72	93.26	8318.
7	7	2015-11-08	0.98	109428.33	703.75	101815.36	80.00	6829.
13	13	2015-09-27	0.99	106803.39	1204.88	99409.21	154.84	6034.
43	43	2015-03-01	0.99	55595.74	629.46	45633.34	181.49	9151.
...
17169	43	2017-03-05	0.99	155011.12	35367.23	5175.81	5.91	114462.
17170	44	2017-02-26	0.99	171145.00	34520.03	6936.39	0.00	129688.
17536	39	2017-04-02	0.98	402676.23	34093.33	58330.53	207.85	310044.
17537	40	2017-03-26	0.90	456645.91	36169.35	51398.72	139.55	368938.
17540	43	2017-03-05	0.99	367519.17	61166.48	55123.99	126.80	251101.

2796 rows × 14 columns



Subsetting based on text data

In [26]:

```
# it will print all rows with "type"="organic"
avocado[avocado["type"]=="organic"]
```

Out[26]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bags
9126	0	2015-12-27	1.83	989.55	8.16	88.59	0.00	892.80
9127	1	2015-12-20	1.89	1163.03	30.24	172.14	0.00	960.65
9128	2	2015-12-13	1.85	995.96	10.44	178.70	0.00	806.82
9129	3	2015-12-06	1.84	1158.42	90.29	104.18	0.00	963.95
9130	4	2015-11-29	1.94	831.69	0.00	94.73	0.00	736.96
...
18244	7	2018-02-04	1.63	17074.83	2046.96	1529.20	0.00	13498.67
18245	8	2018-01-28	1.71	13888.04	1191.70	3431.50	0.00	9264.84
18246	9	2018-01-21	1.87	13766.76	1191.92	2452.79	727.94	9394.11
18247	10	2018-01-14	1.93	16205.22	1527.63	2981.04	727.01	10969.54
18248	11	2018-01-07	1.62	17489.58	2894.77	2356.13	224.53	12014.15

9123 rows × 14 columns



Subsetting based on dates

In [27]:

```
# it will print all rows with "Date" <= 2015-02-04
avocado[avocado["Date"]<="2015-02-04"]
```

Out[27]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bags
47	47	2015-02-01	0.99	70873.60	1353.90	60017.20	179.32	9323.18
48	48	2015-01-25	1.06	45147.50	941.38	33196.16	164.14	10845.82
49	49	2015-01-18	1.17	44511.28	914.14	31540.32	135.77	11921.05
50	50	2015-01-11	1.24	41195.08	1002.85	31640.34	127.12	8424.77
51	51	2015-01-04	1.22	40873.28	2819.50	28287.42	49.90	9716.46
...
11928	46	2015-02-01	1.77	7210.19	1634.42	3012.44	0.00	2563.33
11929	47	2015-01-25	1.63	7324.06	1934.46	3032.72	0.00	2356.88
11930	48	2015-01-18	1.71	5508.20	1793.64	2078.72	0.00	1635.84
11931	49	2015-01-11	1.69	6861.73	1822.28	2377.54	0.00	2661.91
11932	50	2015-01-04	1.64	6182.81	1561.30	2958.17	0.00	1663.34

540 rows × 14 columns



Subsetting using .isin()

In [30]:

```
# subset the avocado in the region Boston or SanDiego
regionFilter = avocado["region"].isin(["Boston", "SanDiego"])
avocado[regionFilter]
```

Out[30]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Ba
208	0	2015-12-27	1.13	450816.39	3886.27	346964.70	13952.56	86012.
209	1	2015-12-20	1.07	489802.88	4912.37	390100.99	5887.72	88901.
210	2	2015-12-13	1.01	549945.76	4641.02	455362.38	219.40	89722.
211	3	2015-12-06	1.02	488679.31	5126.32	407520.22	142.99	75889.
212	4	2015-11-29	1.19	350559.81	3609.25	272719.08	105.86	74125.
...
18100	7	2018-02-04	1.81	17454.74	1158.41	7388.27	0.00	8908.
18101	8	2018-01-28	1.91	17579.47	1145.64	8284.41	0.00	8149.
18102	9	2018-01-21	1.95	18676.37	1088.49	9282.37	0.00	8305.
18103	10	2018-01-14	1.81	21770.02	3285.98	14338.52	0.00	4145.
18104	11	2018-01-07	2.06	16746.82	5150.82	9366.31	0.00	2229.

676 rows × 14 columns



Multiple parameter Filtering

Use logical operators to combine different filters

In [32]: `print(avocado["year"].dtype)`

int64

In [34]: `# subset the avocado in the region Boston or SanDiego in the year 2016 or 2017`
`regionFilter = avocado["region"].isin(["Boston", "SanDiego"])`
`yearFilter = avocado["year"].isin([2016, 2017])`
`avocado[regionFilter & yearFilter]`

Out[34]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bag
3016	0	2016-12-25	1.28	447600.75	4349.63	346516.32	4183.69	92551.1
3017	1	2016-12-18	1.09	579577.33	6123.84	488107.01	7765.43	77581.0
3018	2	2016-12-11	1.22	510800.58	3711.20	409645.98	5052.84	92390.5
3019	3	2016-12-04	1.26	473428.36	4371.95	393748.18	3449.16	71859.0
3020	4	2016-11-27	1.45	391257.01	4243.20	317090.39	3069.37	66854.0
...
16962	48	2017-01-29	1.21	18191.46	1477.75	8949.53	4.86	7759.3
16963	49	2017-01-22	1.73	10842.77	2019.23	6869.87	0.00	1953.6
16964	50	2017-01-15	1.82	11578.42	2529.20	7637.66	0.00	1411.5
16965	51	2017-01-08	1.52	16775.97	2363.28	9429.06	0.00	4983.6
16966	52	2017-01-01	1.45	15752.25	1385.18	8618.28	0.00	5748.7

420 rows × 14 columns



Detecting missing values .isna()

In [35]:

`avocado.isna()`

Out[35]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bags	Small Bags	Large Bags
0	False	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False
...
18244	False	False	False	False	False	False	False	False	False	False
18245	False	False	False	False	False	False	False	False	False	False
18246	False	False	False	False	False	False	False	False	False	False
18247	False	False	False	False	False	False	False	False	False	False
18248	False	False	False	False	False	False	False	False	False	False

18249 rows × 14 columns



We can use .any() function to get a concise info

In [36]: `avocado.isna().any()`

Out[36]:

Unnamed: 0	False
Date	False
AveragePrice	False
Total Volume	False
4046	False
4225	False
4770	False
Total Bags	False
Small Bags	False
Large Bags	False
XLarge Bags	False
type	False
year	False
region	False
dtype: bool	

Counting missing values

In [37]: `avocado.isna().sum()`

```
Out[37]: Unnamed: 0      0
Date          0
AveragePrice  0
Total Volume  0
4046          0
4225          0
4770          0
Total Bags    0
Small Bags    0
Large Bags    0
XLarge Bags   0
type          0
year          0
region        0
dtype: int64
```

Removing missing values

```
In [38]: # Luckily we don't have any NaN but if we have we can use any of the two
avocado.dropna()

# **** OR ****

meanVal = avocado["AveragePrice"].mean()
avocado.fillna(meanVal)
```

Out[38]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bags
0	0	2015-12-27	1.33	64236.62	1036.74	54454.85	48.16	8696.87
1	1	2015-12-20	1.35	54876.98	674.28	44638.81	58.33	9505.56
2	2	2015-12-13	0.93	118220.22	794.70	109149.67	130.50	8145.35
3	3	2015-12-06	1.08	78992.15	1132.00	71976.41	72.58	5811.16
4	4	2015-11-29	1.28	51039.60	941.48	43838.39	75.78	6183.95
...
18244	7	2018-02-04	1.63	17074.83	2046.96	1529.20	0.00	13498.67
18245	8	2018-01-28	1.71	13888.04	1191.70	3431.50	0.00	9264.84
18246	9	2018-01-21	1.87	13766.76	1191.92	2452.79	727.94	9394.11
18247	10	2018-01-14	1.93	16205.22	1527.63	2981.04	727.01	10969.54
18248	11	2018-01-07	1.62	17489.58	2894.77	2356.13	224.53	12014.15

18249 rows × 14 columns



Adding a new column

In [39]:

```
avocado["AveragePricePer100"] = avocado["AveragePrice"]*100
avocado
```

Out[39]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bags
0	0	2015-12-27	1.33	64236.62	1036.74	54454.85	48.16	8696.87
1	1	2015-12-20	1.35	54876.98	674.28	44638.81	58.33	9505.56
2	2	2015-12-13	0.93	118220.22	794.70	109149.67	130.50	8145.35
3	3	2015-12-06	1.08	78992.15	1132.00	71976.41	72.58	5811.16
4	4	2015-11-29	1.28	51039.60	941.48	43838.39	75.78	6183.95
...
18244	7	2018-02-04	1.63	17074.83	2046.96	1529.20	0.00	13498.67
18245	8	2018-01-28	1.71	13888.04	1191.70	3431.50	0.00	9264.84
18246	9	2018-01-21	1.87	13766.76	1191.92	2452.79	727.94	9394.11
18247	10	2018-01-14	1.93	16205.22	1527.63	2981.04	727.01	10969.54
18248	11	2018-01-07	1.62	17489.58	2894.77	2356.13	224.53	12014.15

18249 rows × 15 columns



Deleting columns in DataFrame .drop(lst, axis=1)

In [40]:

avocado.drop(["AveragePricePer100"], axis=1)

Out[40]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bags
0	0	2015-12-27	1.33	64236.62	1036.74	54454.85	48.16	8696.87
1	1	2015-12-20	1.35	54876.98	674.28	44638.81	58.33	9505.56
2	2	2015-12-13	0.93	118220.22	794.70	109149.67	130.50	8145.35
3	3	2015-12-06	1.08	78992.15	1132.00	71976.41	72.58	5811.16
4	4	2015-11-29	1.28	51039.60	941.48	43838.39	75.78	6183.95
...
18244	7	2018-02-04	1.63	17074.83	2046.96	1529.20	0.00	13498.67
18245	8	2018-01-28	1.71	13888.04	1191.70	3431.50	0.00	9264.84
18246	9	2018-01-21	1.87	13766.76	1191.92	2452.79	727.94	9394.11
18247	10	2018-01-14	1.93	16205.22	1527.63	2981.04	727.01	10969.54
18248	11	2018-01-07	1.62	17489.58	2894.77	2356.13	224.53	12014.15

18249 rows × 14 columns



Summary statistics

In [41]: # mean of the AveragePrice of avocado
avocado["AveragePrice"].mean()

Out[41]: np.float64(1.405978409775878)

Summarizing dates

In [42]: avocado["Date"].max()

Out[42]: '2018-03-25'

.agg() method

Pandas Series.agg() is used to pass a function or list of function to be applied on a series or even each element of series separately

```
In [44]: def pct30(column):
    #return the 0.3 quartile
    return column.quantile(0.3)
def pct50(column):
    #return the 0.5 quartile
    return column.quantile(0.5)

avocado[["AveragePrice","Total Bags"]].agg([pct30,pct50])
```

```
Out[44]:   AveragePrice  Total Bags
pct30          1.15    7316.634
pct50          1.37   39743.830
```

Dropping duplicate names .drop_duplicates(lst)

```
In [45]: temp = avocado.drop_duplicates(subset=["year"])
temp
```

```
Out[45]:      Unnamed: 0      Date  AveragePrice  Total Volume  4046  4225  4770  Total Bags
0             0  2015-12-27        1.33    64236.62  1036.74  54454.85  48.16  8696.87
2808          0  2016-12-25        1.52    73341.73  3202.39  58280.33  426.92 11432.09
5616          0  2017-12-31        1.47   113514.42  2622.70 101135.53  20.25  9735.94
8478          0  2018-03-25        1.57  149396.50  16361.69 109045.03  65.45 23924.33
```

Count categorical data .value_counts()

```
In [46]: # count number of avocado in each year in descending order
avocado["year"].value_counts(sort=True, ascending = False)
```

```
Out[46]: year
2017    5722
2016    5616
2015    5615
2018    1296
Name: count, dtype: int64
```

Grouped summaries .groupby(col)

This function will group similar categories into one and then we can perform some summary statistics

```
In [51]: # group by multiple columns and perform multiple summary statistic operations
avocado.groupby(["year", "type"])["AveragePrice"].agg(["min", "max", "mean", "median"])
```

Out[51]:

			min	max	mean	median
	year	type				
2015	conventional	0.49	1.59	1.077963	1.08	
		organic	0.81	2.79	1.673324	1.67
2016	conventional	0.51	2.20	1.105595	1.08	
		organic	0.58	3.25	1.571684	1.53
2017	conventional	0.46	2.22	1.294888	1.30	
		organic	0.44	3.17	1.735521	1.72
2018	conventional	0.56	1.74	1.127886	1.14	
		organic	1.01	2.30	1.567176	1.55

Pivot table

A pivot table is a table of statistics that summarizes the data of a more extensive table

```
In [50]: # this is the same table we build in the previous cell but using pivot table
avocado.pivot_table(index=["year", "type"], aggfunc=["min", "max", "mean", "median"],
```

Out[50]:

			min	max	mean	median
			AveragePrice	AveragePrice	AveragePrice	AveragePrice
	year	type				
2015	conventional		0.49	1.59	1.077963	1.08
		organic	0.81	2.79	1.673324	1.67
2016	conventional		0.51	2.20	1.105595	1.08
		organic	0.58	3.25	1.571684	1.53
2017	conventional		0.46	2.22	1.294888	1.30
		organic	0.44	3.17	1.735521	1.72
2018	conventional		0.56	1.74	1.127886	1.14
		organic	1.01	2.30	1.567176	1.55

Explicit indexes

Indexes makes subsetting simpler using .loc and .iloc

Setting column as the index

```
In [52]: regionIndex = avocado.set_index(["region"])
regionIndex
```

Out[52]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4
region							
Albany	0	2015-12-27	1.33	64236.62	1036.74	54454.85	48
Albany	1	2015-12-20	1.35	54876.98	674.28	44638.81	58
Albany	2	2015-12-13	0.93	118220.22	794.70	109149.67	130
Albany	3	2015-12-06	1.08	78992.15	1132.00	71976.41	72
Albany	4	2015-11-29	1.28	51039.60	941.48	43838.39	71
...
WestTexNewMexico	7	2018-02-04	1.63	17074.83	2046.96	1529.20	0
WestTexNewMexico	8	2018-01-28	1.71	13888.04	1191.70	3431.50	0
WestTexNewMexico	9	2018-01-21	1.87	13766.76	1191.92	2452.79	72
WestTexNewMexico	10	2018-01-14	1.93	16205.22	1527.63	2981.04	72
WestTexNewMexico	11	2018-01-07	1.62	17489.58	2894.77	2356.13	224

18249 rows × 14 columns

In [53]:

```
#Instead of doing this
avocado[avocado["region"].isin(["Albany","WestTexNewMexico"])]
```

Out[53]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bags
0	0	2015-12-27	1.33	64236.62	1036.74	54454.85	48.16	8696.87
1	1	2015-12-20	1.35	54876.98	674.28	44638.81	58.33	9505.56
2	2	2015-12-13	0.93	118220.22	794.70	109149.67	130.50	8145.35
3	3	2015-12-06	1.08	78992.15	1132.00	71976.41	72.58	5811.16
4	4	2015-11-29	1.28	51039.60	941.48	43838.39	75.78	6183.95
...
18244	7	2018-02-04	1.63	17074.83	2046.96	1529.20	0.00	13498.67
18245	8	2018-01-28	1.71	13888.04	1191.70	3431.50	0.00	9264.84
18246	9	2018-01-21	1.87	13766.76	1191.92	2452.79	727.94	9394.11
18247	10	2018-01-14	1.93	16205.22	1527.63	2981.04	727.01	10969.54
18248	11	2018-01-07	1.62	17489.58	2894.77	2356.13	224.53	12014.15

673 rows × 15 columns



In [55]:

```
# we can simply do
regionIndex.loc[["Albany", "WestTexNewMexico"]]
```

Out[55]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4
region							
Albany	0	2015-12-27	1.33	64236.62	1036.74	54454.85	48
Albany	1	2015-12-20	1.35	54876.98	674.28	44638.81	58
Albany	2	2015-12-13	0.93	118220.22	794.70	109149.67	130
Albany	3	2015-12-06	1.08	78992.15	1132.00	71976.41	72
Albany	4	2015-11-29	1.28	51039.60	941.48	43838.39	71
...
WestTexNewMexico	7	2018-02-04	1.63	17074.83	2046.96	1529.20	0
WestTexNewMexico	8	2018-01-28	1.71	13888.04	1191.70	3431.50	0
WestTexNewMexico	9	2018-01-21	1.87	13766.76	1191.92	2452.79	72
WestTexNewMexico	10	2018-01-14	1.93	16205.22	1527.63	2981.04	72
WestTexNewMexico	11	2018-01-07	1.62	17489.58	2894.77	2356.13	224

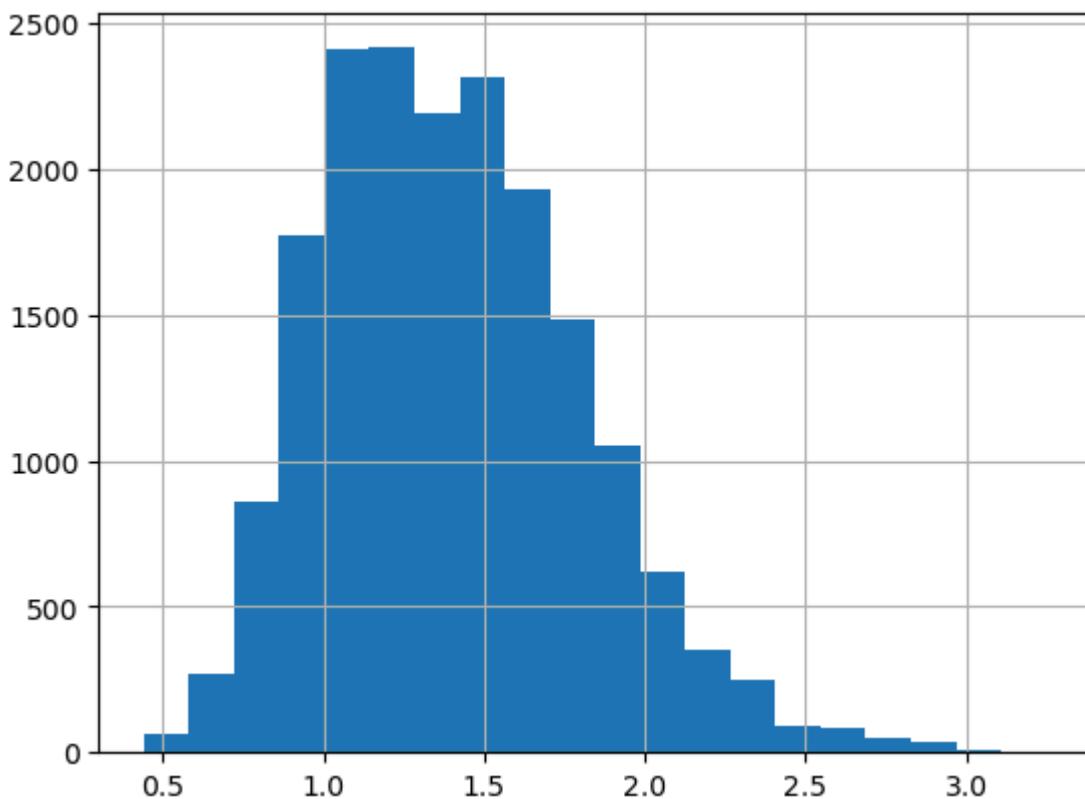
673 rows × 14 columns



Visualizing your data

In [57]:

```
#Histograms
avocado["AveragePrice"].hist(bins=20)
plt.show()
```

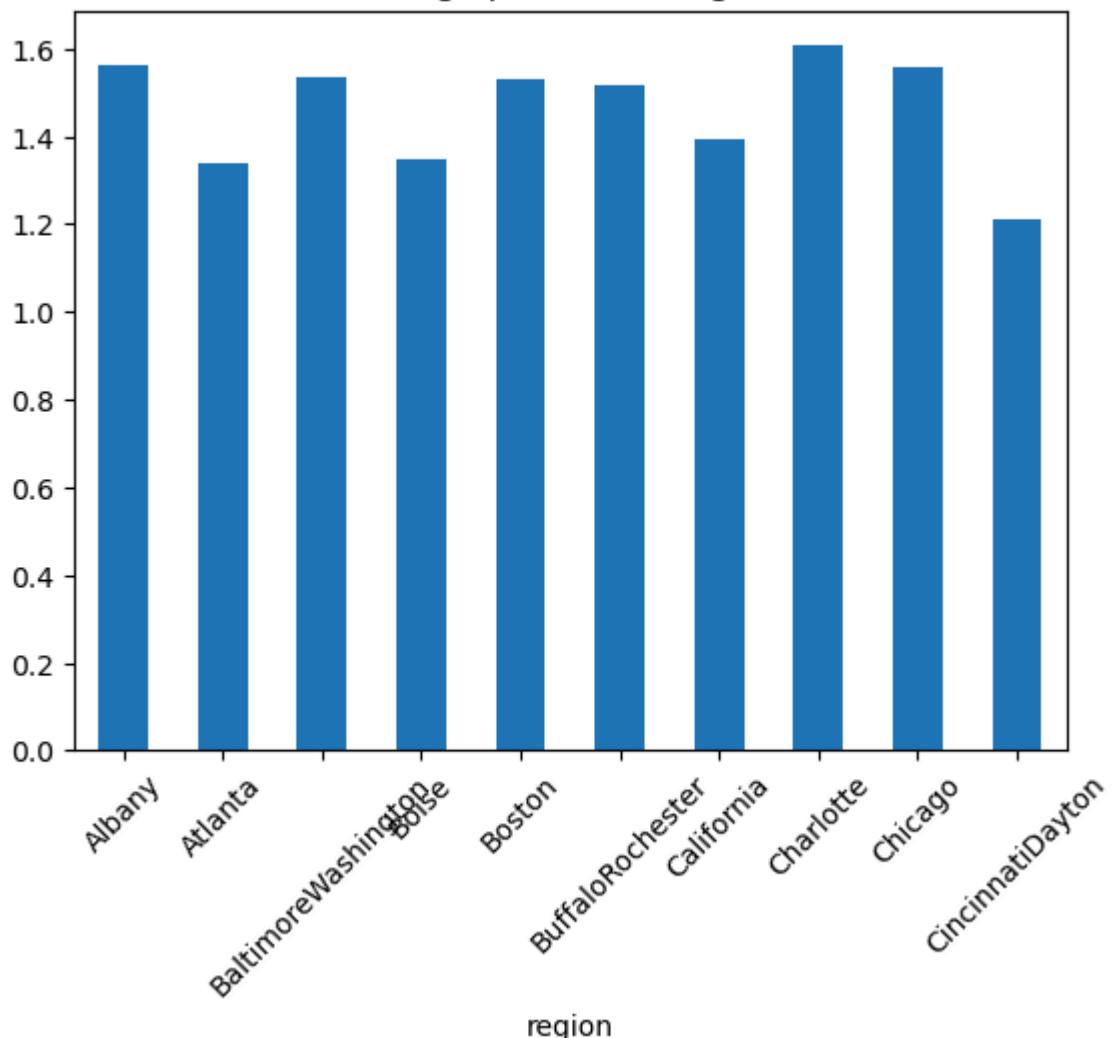


```
In [58]: #Bar plots
regionFilter = avocado.groupby("region")["AveragePrice"].mean().head(10)
regionFilter
```

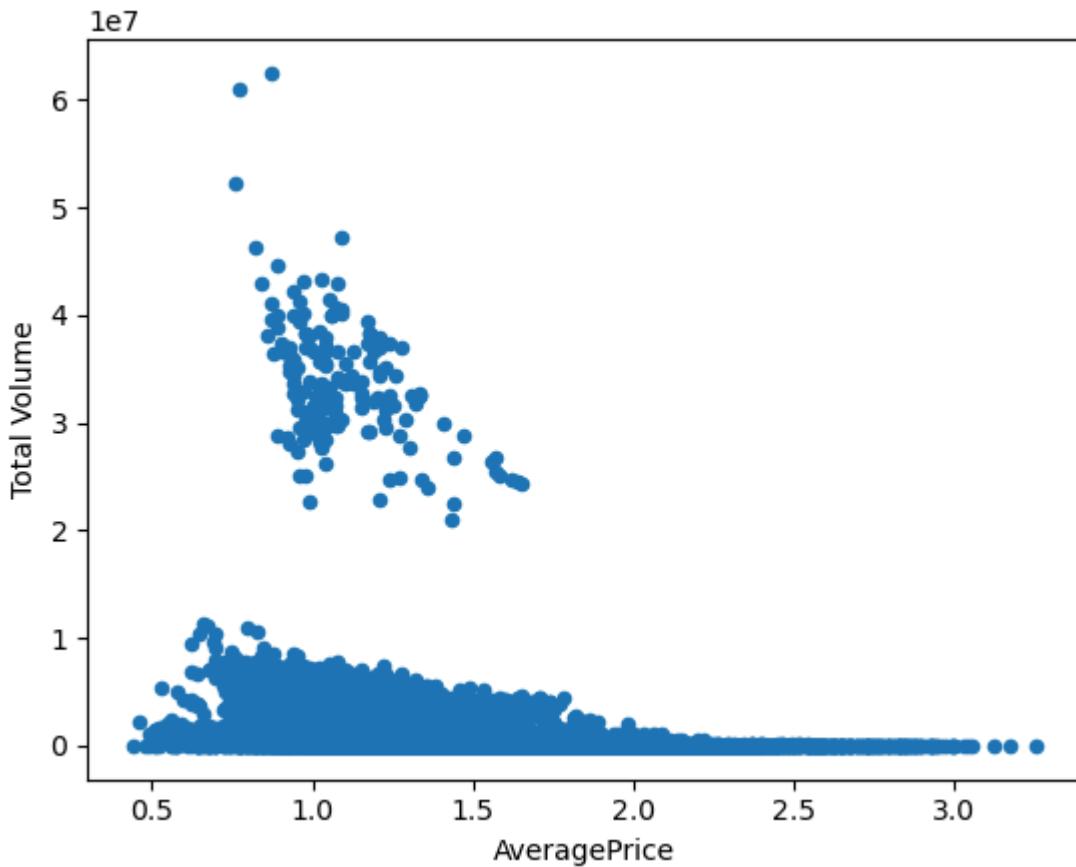
```
Out[58]: region
Albany           1.561036
Atlanta          1.337959
BaltimoreWashington 1.534231
Boise            1.348136
Boston            1.530888
BuffaloRochester 1.516834
California        1.395325
Charlotte          1.606036
Chicago             1.556775
CincinnatiDayton  1.209201
Name: AveragePrice, dtype: float64
```

```
In [60]: regionFilter.plot(kind = "bar", rot=45, title="Average price in 10 regions")
plt.show()
```

Average price in 10 regions



```
In [63]: # Scatter plot
avocado.plot(x="AveragePrice",y="Total Volume",kind="scatter")
plt.show()
```



Arithmetic with Series & DataFrames

```
In [64]: # subtract Averages with AveragePriceee : P  
# Dah its 0  
avocado[ "AveragePrice" ].sub(avocado[ "AveragePrice" ])
```

```
Out[64]: 0      0.0  
1      0.0  
2      0.0  
3      0.0  
4      0.0  
...  
18244  0.0  
18245  0.0  
18246  0.0  
18247  0.0  
18248  0.0  
Name: AveragePrice, Length: 18249, dtype: float64
```

```
In [ ]:
```